

## Training the data on CNN

In [1]:

```
from sklearn.model_selection import KFold, StratifiedKFold
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.models import Sequential, load_model
from tensorflow.keras.layers import Conv2D, Dropout, RandomFlip, RandomRotation, RandomZoom, RandomContrast, RandomCrop,
from tensorflow.keras.losses import SparseCategoricalCrossentropy
from tensorflow.keras.optimizers import SGD, Adam
from sklearn.model_selection import train_test_split, GridSearchCV
from keras.wrappers.scikit_learn import KerasClassifier
from sklearn.metrics import accuracy_score, confusion_matrix
from tensorflow.keras.layers import BatchNormalization
from numpy import mean, std
from matplotlib import pyplot as plt
import tensorflow as tf
import numpy as np
import pandas as pd
import numpy.random as npr
import matplotlib.pyplot as plt
%matplotlib inline
plt.style.use('bmh')

# Loading Data
data_train = np.load('data_train.npy')
labels_train = np.load('new_labels_train.npy')

print(data_train.shape, labels_train.shape)

# Counting number samples per class
vals, counts = np.unique(labels_train, return_counts=True)

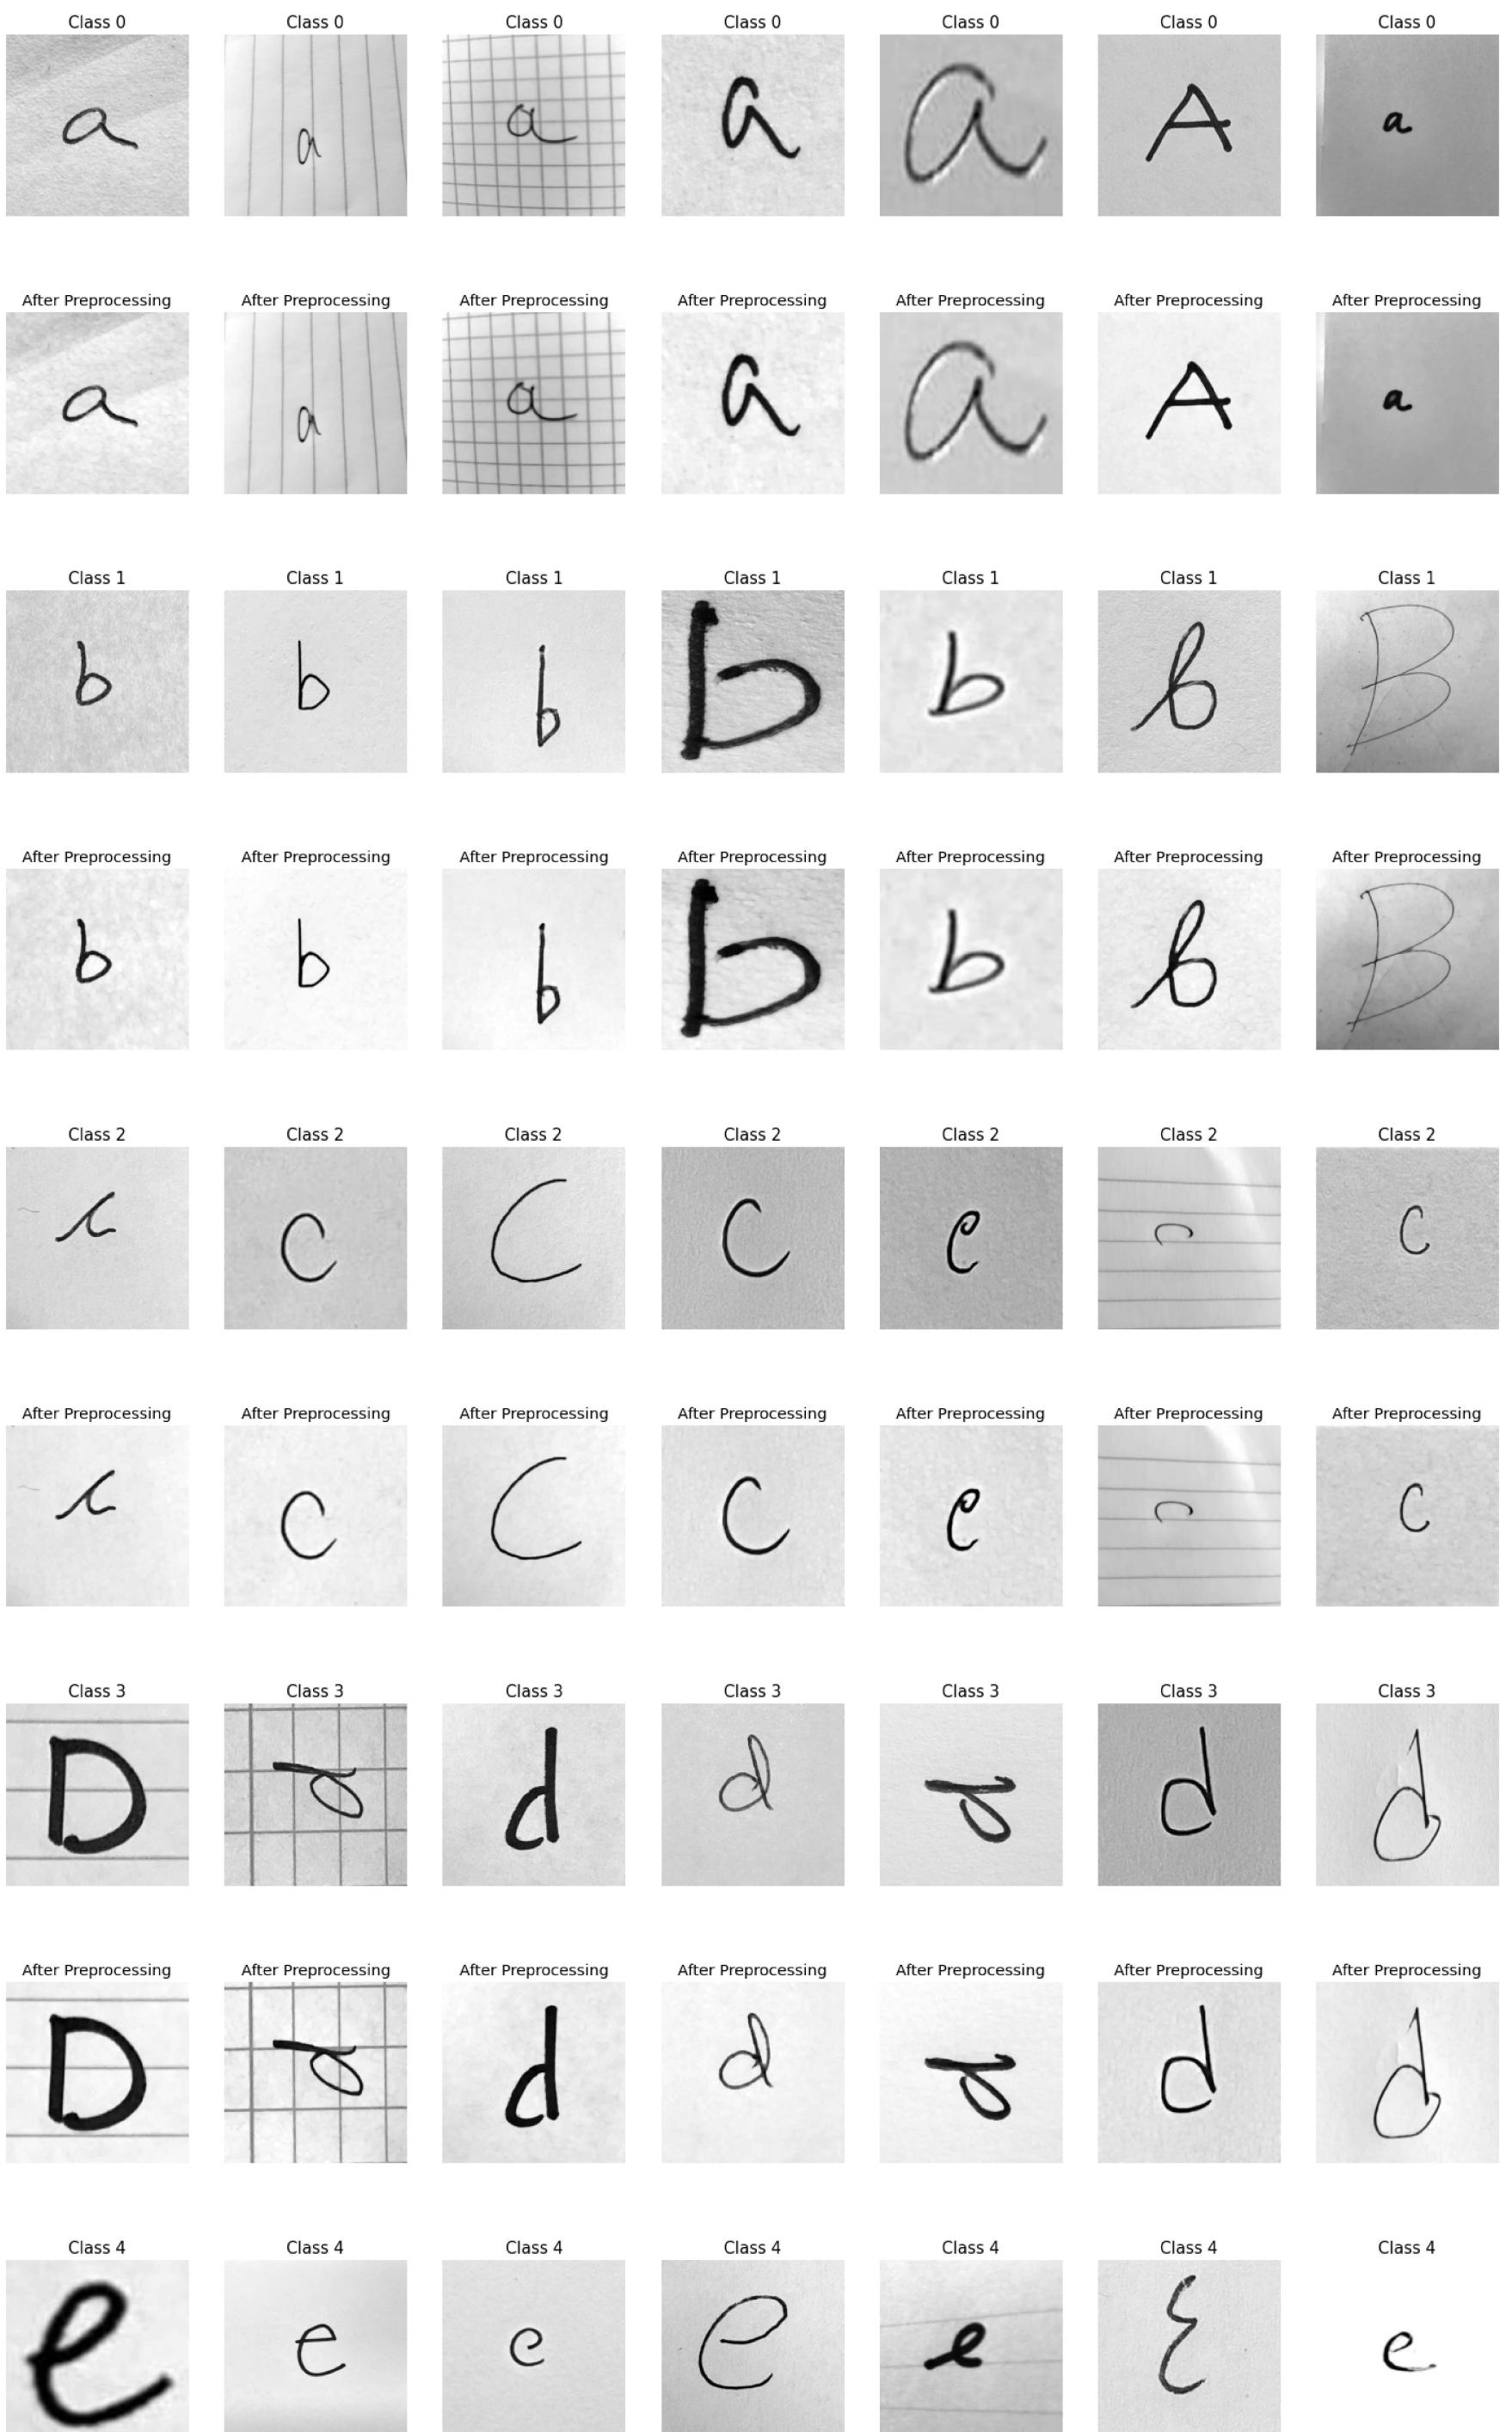
plt.bar(vals, counts)
plt.xticks(range(10), range(10))
plt.xlabel('Classes', size=20)
plt.ylabel('# Samples per Class', size=20)
plt.title('Training Data (Total = '+str(data_train.shape[1])+' samples)', size=15);
```

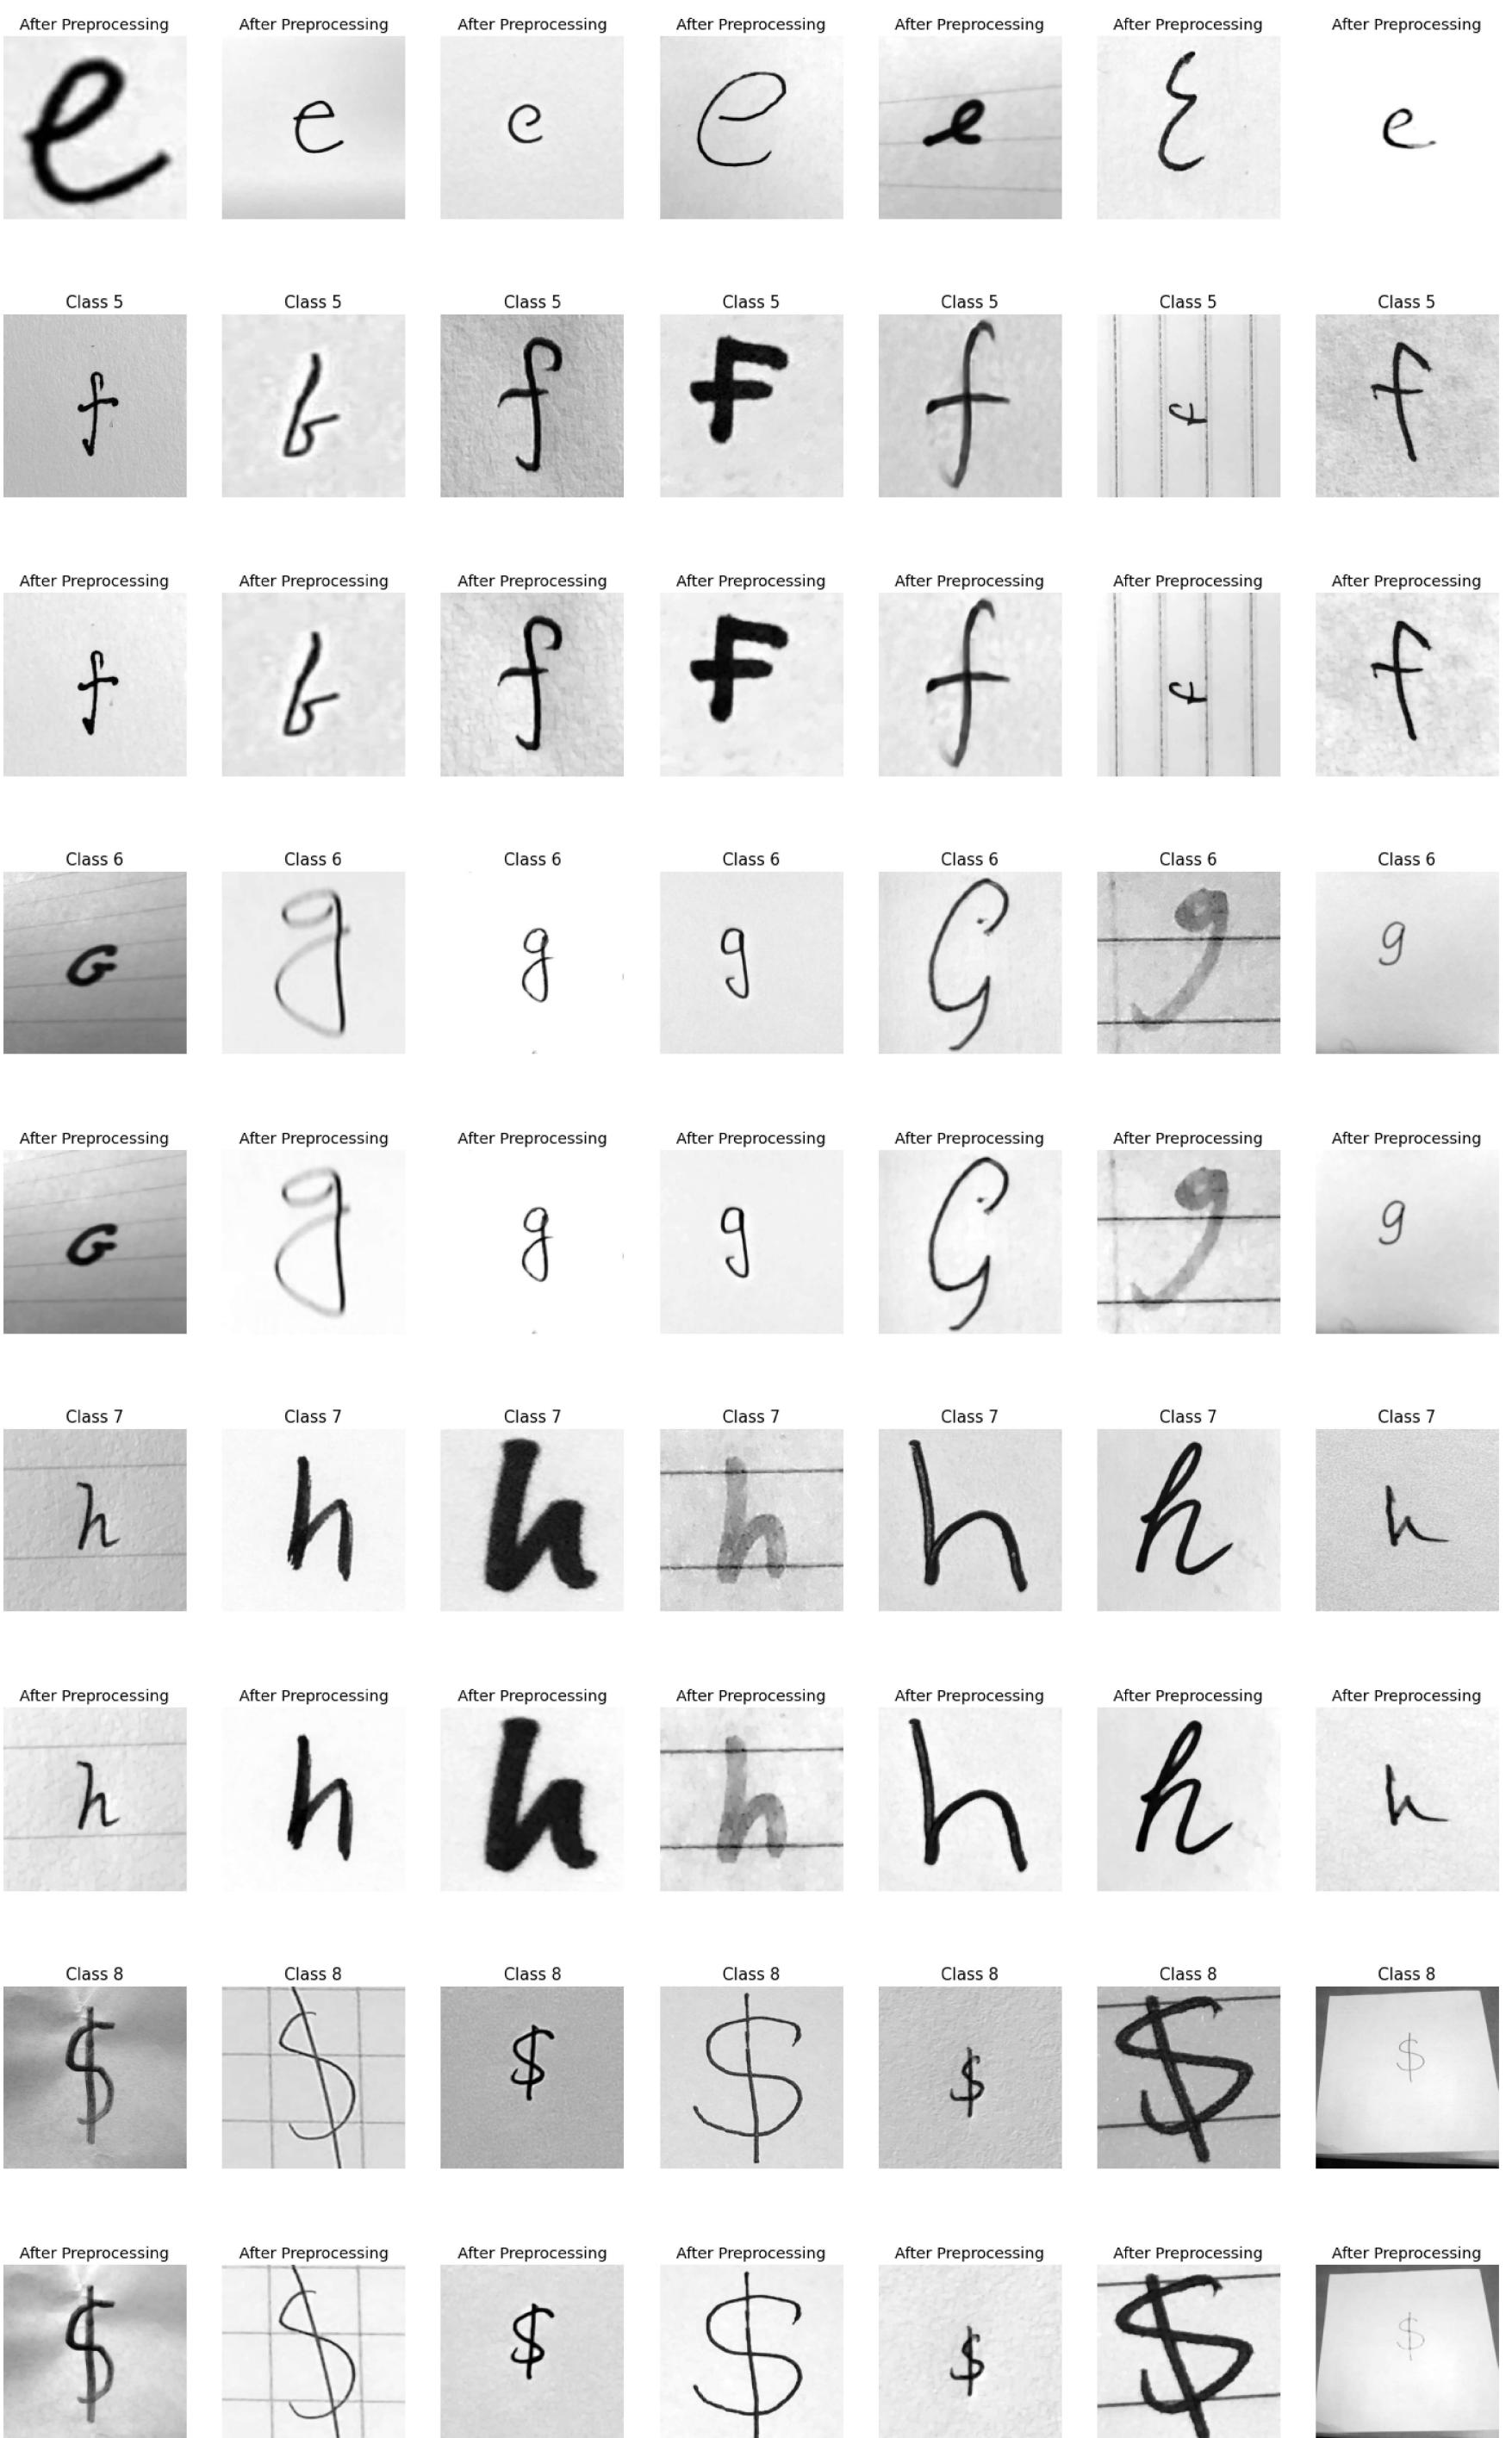


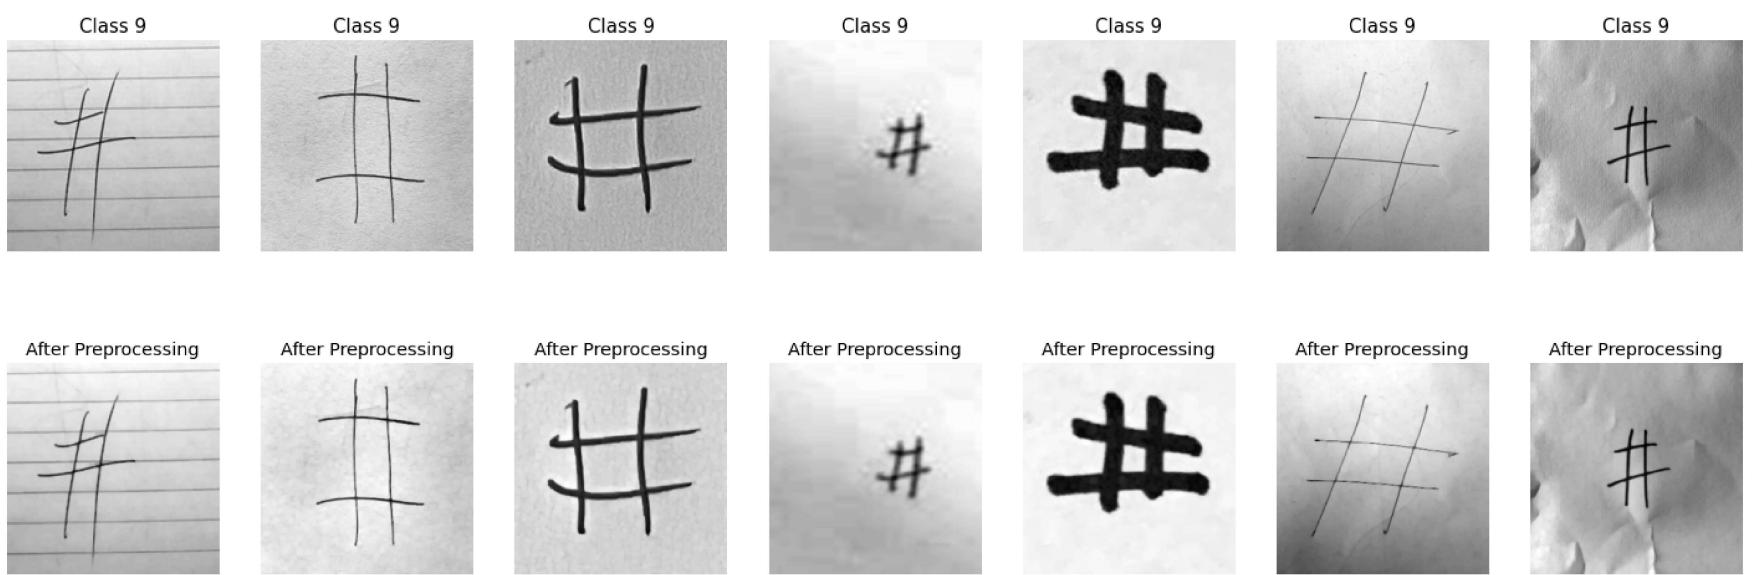
## Visualising some images after and before preprocessing

In [2]:

```
import cv2
kernel = np.ones((5,5),np.uint8)
for i in range(0,10):
    rnd_sample = npr.permutation(np.where(labels_train==i)[0])
    fig1=plt.figure(figsize=(25,20))
    for j in range(7):
        fig1.add_subplot(5,7,j+1)
        plt.imshow(data_train[:,rnd_sample[j]].reshape((300,300)),cmap='gray')
        plt.axis('off');plt.title('Class '+str(int(labels_train[rnd_sample[j]])),size=15)
    plt.show()
    print('\n\n')
    j = 0
    fig=plt.figure(figsize=(25,20))
    dt = []
    k=0
    for j in range(7):
        fig.add_subplot(5,7,j+1)
        dt.append(cv2.medianBlur((data_train[:,rnd_sample[j]].reshape((300,300))), 3))
        dt[j] = cv2.morphologyEx(dt[j], cv2.MORPH_OPEN, kernel)
        plt.imshow(dt[j],cmap='gray')
        k = k + 1
        plt.axis('off');plt.title('After Preprocessing')
    plt.show()
    print('\n\n')
```







In [3]:

```
# import argparse
# import skimage as sk
# from skimage import transform
# from skimage import util
import cv2
train_data_length = data_train.shape[1]
# Initial function for Loading training set and preprocess dataset
def load_dataset():
    # Load dataset
    data_rgb = []
    kernel = np.ones((4,4),np.uint8)
    for i in range(train_data_length):
        data_rgb.append(data_train[:,i].reshape(300,300))
        data_rgb[i] = cv2.medianBlur(data_rgb[i], 3)
        data_rgb[i] = cv2.morphologyEx(data_rgb[i], cv2.MORPH_OPEN, kernel)
        data_rgb[i] = cv2.resize(data_rgb[i], (50,50), interpolation=cv2.INTER_AREA)
    data_rgb = np.array(data_rgb)
    print(data_rgb.shape)

    # Using 80-20 split
    trainX, testX, trainY, testY = train_test_split(data_rgb, labels_train, test_size = 0.2, stratify=labels_train)

    # reshape dataset to have a single channel
    trainX = trainX.reshape((trainX.shape[0], 50,50, 1))
    testX = testX.reshape((testX.shape[0], 50,50, 1))

    # one hot encode target values
    trainY = to_categorical(trainY)
    testY = to_categorical(testY)
    print(trainX.shape, trainY.shape, testX.shape, testY.shape)
    return trainX, trainY, testX, testY
```

In [4]:

```
# scale pixels
def prep_pixels(train, test):
    # convert from integers to floats
    train_norm = train.astype('float32')
    test_norm = test.astype('float32')

    # he_uniformize to range 0-1
    train_norm = (train_norm) / 255.0
    test_norm = (test_norm) / 255.0

    # return he_uniformized images
    return train_norm, test_norm
```

In [5]:

```
# define cnn model
def define_model(learn_rate = '0.007'):
    model = Sequential()
    model.add(RandomContrast((0,0.3),input_shape=(50,50,1)))
    model.add(RandomRotation(0.1, fill_mode='reflect', interpolation='bilinear'))
    model.add(Conv2D(32, (5, 5), activation='relu', kernel_initializer='he_uniform', input_shape=(50,50, 1)))
    model.add(Conv2D(32, (5, 5), activation='relu', kernel_initializer='he_uniform'))
    model.add(BatchNormalization())
    model.add(MaxPooling2D((2, 2)))
    model.add(Dropout(0.5))
    model.add(Conv2D(32, (3,3),activation='relu', kernel_initializer='he_uniform'))
    model.add(Conv2D(64, (3,3),activation='relu', kernel_initializer='he_uniform'))
    model.add(BatchNormalization())
    model.add(MaxPooling2D((2, 2)))
    model.add(Dropout(0.5))
    model.add(Conv2D(64,(3,3), activation='relu', kernel_initializer='he_uniform'))
    model.add(Conv2D(64,(3,3), activation='relu', kernel_initializer='he_uniform'))
    model.add(BatchNormalization())
    model.add(MaxPooling2D((2, 2)))
    model.add(Dropout(0.5))
    model.add(Flatten())
    model.add(Dense(256, activation='relu', kernel_initializer='he_uniform'))
    model.add(Dropout(0.4))
    model.add(Dense(128, activation='relu', kernel_initializer='he_uniform'))
    model.add(Dropout(0.4))
    model.add(Dense(10, activation='softmax'))
```

```
# compile model
opt = Adam(learning_rate = learn_rate , epsilon=1e-5)
model.compile(optimizer=opt, loss='categorical_crossentropy', metrics=['accuracy'])
return model
```

In [6]:

```
# Function for Finding best parameters using GRID CV
def training(dataX, dataY, tes_X, tes_y):
    parameters = {'epochs':[800], 'batch_size':[32, 128], 'learn_rate':[0.001, 0.0007]}
    model = KerasClassifier(build_fn = define_model, epochs = 800, batch_size = 32)
    grid = GridSearchCV(estimator = model, param_grid = parameters, n_jobs = 1, cv = 3, verbose = 2, return_train_score =
grid_result = grid.fit(dataX, dataY)

    # summarize results
    print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))
    means = grid_result.cv_results_['mean_test_score']
    stds = grid_result.cv_results_['std_test_score']
    params = grid_result.cv_results_['params']
    for mean, stdev, param in zip(means, stds, params):
        print("%f (%f) with: %r" % (mean, stdev, param))
    return grid_result.best_estimator_
```

**Code for GRID Search and Best estimator in comments below, you don't have to run these 2 cells below as we know our best parameters now.**

In [13]:

```
# def Load_and_train():
#     # Load dataset
#     trainX, trainY, testX, testY = Load_dataset()

#     # prepare pixel data
#     trainX, testX = prep_pixels(trainX, testX)

#     # evaluate model
#     best_estimator = training(trainX, trainY, testX, testY)
#     return best_estimator
# best_estimator = Load_and_train()
```

In [10]:

```
# print(best_estimator.get_params())
{'epochs': 800, 'batch_size': 32, 'learn_rate': 0.0007, 'build_fn': <function define_model at 0x2b6665dff5e0>}
```

**We can see from above what the best parameters are, so we will train our model now on these parameters.**

In [7]:

```
def final_training(dataX, dataY, testX, testY):
    model = define_model(learn_rate = 0.0007)
    print(model.summary())
    trainX, trainY, testX, testY = dataX, dataY, testX, testY
    # fit model
    history = model.fit(trainX, trainY, epochs = 800, batch_size = 32, validation_data = (testX, testY), verbose=1)
    print("Model Fitted")
    return model, history
```

In [8]:

```
# plot diagnostic learning curves
def summarize_diagnostics(histories):
    plt.figure(figsize=(12,7))
    plt.plot(history.history['accuracy'])
    plt.plot(history.history['val_accuracy'])
    plt.title('Training and validation Accuracy')
    plt.ylabel('Accuracy')
    plt.xlabel('epoch')
    plt.legend(['Accuracy', 'Val Accuracy'], loc='lower right')
    plt.show()
    plt.figure(figsize=(12,7))
    plt.plot(history.history['loss'])
    plt.plot(history.history['val_loss'])
    plt.title('Training and validation Loss')
    plt.ylabel('Loss')
    plt.xlabel('Epochs')
    plt.legend(['Loss', 'Val Loss'], loc='upper right')
    plt.show()
```

**This below cell is where we call 3 functions to split and preprocess the dataset, Normalize the images and finally train on 80% data.**

In [9]:

```
trainX, trainY, testX, testY = load_dataset()

# prepare pixel data
trainX, testX = prep_pixels(trainX, testX)

# training the model
model, history = final_training(trainX, trainY, testX, testY)
```

(6720, 50, 50)

(5376, 50, 50, 1) (5376, 10) (1344, 50, 50, 1) (1344, 10)  
2022-04-20 13:56:29.748779: I tensorflow/core/platform/cpu\_feature\_guard.cc:142] This TensorFlow binary is optimized with oneAPI Deep Neural Network Library (oneDNN) to use the following CPU instructions in performance-critical operations: SS E4.1 SSE4.2 AVX AVX2 FMA  
To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.  
2022-04-20 13:56:30.252276: I tensorflow/core/common\_runtime/gpu/gpu\_device.cc:1510] Created device /job:localhost/replica:0/task:0/device:GPU:0 with 79125 MB memory: -> device: 0, name: NVIDIA A100-SXM4-80GB, pci bus id: 0000:47:00.0, compute capability: 8.0  
Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
random_contrast (RandomContrast)	(None, 50, 50, 1)	0
random_rotation (RandomRotation)	(None, 50, 50, 1)	0
conv2d (Conv2D)	(None, 46, 46, 32)	832
conv2d_1 (Conv2D)	(None, 42, 42, 32)	25632
batch_normalization (BatchNormal)	(None, 42, 42, 32)	128
max_pooling2d (MaxPooling2D)	(None, 21, 21, 32)	0
dropout (Dropout)	(None, 21, 21, 32)	0
conv2d_2 (Conv2D)	(None, 19, 19, 32)	9248
conv2d_3 (Conv2D)	(None, 17, 17, 64)	18496
batch_normalization_1 (BatchNormal)	(None, 17, 17, 64)	256
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 64)	0
dropout_1 (Dropout)	(None, 8, 8, 64)	0
conv2d_4 (Conv2D)	(None, 6, 6, 64)	36928
conv2d_5 (Conv2D)	(None, 4, 4, 64)	36928
batch_normalization_2 (BatchNormal)	(None, 4, 4, 64)	256
max_pooling2d_2 (MaxPooling2D)	(None, 2, 2, 64)	0
dropout_2 (Dropout)	(None, 2, 2, 64)	0
flatten (Flatten)	(None, 256)	0
dense (Dense)	(None, 256)	65792
dropout_3 (Dropout)	(None, 256)	0
dense_1 (Dense)	(None, 128)	32896
dropout_4 (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 10)	1290
=====		
Total params:	228,682	
Trainable params:	228,362	
Non-trainable params:	320	

None  
2022-04-20 13:56:31.266568: I tensorflow/compiler/mlir/mlir\_graph\_optimization\_pass.cc:185] None of the MLIR Optimization Passes are enabled (registered 2)  
Epoch 1/800  
2022-04-20 13:56:33.352678: I tensorflow/stream\_executor/cuda/cuda\_dnn.cc:369] Loaded cuDNN version 8201  
2022-04-20 13:56:34.264488: W tensorflow/stream\_executor/gpu/asm\_compiler.cc:77] Couldn't get ptxas version string: Internal: Running ptxas --version returned 32512  
2022-04-20 13:56:34.369035: W tensorflow/stream\_executor/gpu/redzone\_allocator.cc:314] Internal: ptxas exited with non-zero error code 32512, output:  
Relying on driver to perform ptx compilation.  
Modify \$PATH to customize ptxas location.  
This message will be only logged once.  
2022-04-20 13:56:35.782709: I tensorflow/stream\_executor/cuda/cuda\_blas.cc:1760] TensorFloat-32 will be used for the matrix multiplication. This will only be logged once.  
168/168 [=====] - 6s 6ms/step - loss: 3.4276 - accuracy: 0.1025 - val\_loss: 2.3224 - val\_accuracy: 0.1079  
Epoch 2/800  
168/168 [=====] - 1s 5ms/step - loss: 2.4370 - accuracy: 0.1256 - val\_loss: 2.2919 - val\_accuracy: 0.1287  
Epoch 3/800  
168/168 [=====] - 1s 4ms/step - loss: 2.3497 - accuracy: 0.1198 - val\_loss: 2.2682 - val\_accuracy: 0.1659  
Epoch 4/800  
168/168 [=====] - 1s 5ms/step - loss: 2.2975 - accuracy: 0.1295 - val\_loss: 2.2488 - val\_accuracy: 0.1696  
Epoch 5/800  
168/168 [=====] - 1s 5ms/step - loss: 2.2669 - accuracy: 0.1479 - val\_loss: 2.1815 - val\_accuracy: 0.1927  
Epoch 6/800  
168/168 [=====] - 1s 5ms/step - loss: 2.2229 - accuracy: 0.1788 - val\_loss: 2.1117 - val\_accuracy: 0.2277  
Epoch 7/800

168/168 [=====] - 1s 4ms/step - loss: 2.1738 - accuracy: 0.1970 - val\_loss: 2.0937 - val\_accuracy: 0.2455  
Epoch 8/800  
168/168 [=====] - 1s 5ms/step - loss: 2.1528 - accuracy: 0.2052 - val\_loss: 2.0337 - val\_accuracy: 0.2448  
Epoch 9/800  
168/168 [=====] - 1s 5ms/step - loss: 2.1084 - accuracy: 0.2108 - val\_loss: 1.9826 - val\_accuracy: 0.2805  
Epoch 10/800  
168/168 [=====] - 1s 4ms/step - loss: 2.0672 - accuracy: 0.2193 - val\_loss: 1.8753 - val\_accuracy: 0.3103  
Epoch 11/800  
168/168 [=====] - 1s 4ms/step - loss: 2.0255 - accuracy: 0.2461 - val\_loss: 1.9378 - val\_accuracy: 0.2969  
Epoch 12/800  
168/168 [=====] - 1s 4ms/step - loss: 1.9790 - accuracy: 0.2675 - val\_loss: 1.8436 - val\_accuracy: 0.3356  
Epoch 13/800  
168/168 [=====] - 1s 5ms/step - loss: 1.9527 - accuracy: 0.2673 - val\_loss: 1.7890 - val\_accuracy: 0.3497  
Epoch 14/800  
168/168 [=====] - 1s 5ms/step - loss: 1.9237 - accuracy: 0.2811 - val\_loss: 1.6658 - val\_accuracy: 0.3981  
Epoch 15/800  
168/168 [=====] - 1s 5ms/step - loss: 1.8854 - accuracy: 0.3136 - val\_loss: 1.6574 - val\_accuracy: 0.4204  
Epoch 16/800  
168/168 [=====] - 1s 4ms/step - loss: 1.8559 - accuracy: 0.3203 - val\_loss: 1.5604 - val\_accuracy: 0.4494  
Epoch 17/800  
168/168 [=====] - 1s 5ms/step - loss: 1.8001 - accuracy: 0.3432 - val\_loss: 1.5853 - val\_accuracy: 0.4397  
Epoch 18/800  
168/168 [=====] - 1s 5ms/step - loss: 1.7185 - accuracy: 0.3823 - val\_loss: 1.3914 - val\_accuracy: 0.5685  
Epoch 19/800  
168/168 [=====] - 1s 5ms/step - loss: 1.6457 - accuracy: 0.4131 - val\_loss: 1.4740 - val\_accuracy: 0.4635  
Epoch 20/800  
168/168 [=====] - 1s 4ms/step - loss: 1.5861 - accuracy: 0.4306 - val\_loss: 1.4350 - val\_accuracy: 0.4874  
Epoch 21/800  
168/168 [=====] - 1s 4ms/step - loss: 1.5362 - accuracy: 0.4554 - val\_loss: 1.3044 - val\_accuracy: 0.5655  
Epoch 22/800  
168/168 [=====] - 1s 5ms/step - loss: 1.4792 - accuracy: 0.4829 - val\_loss: 1.4089 - val\_accuracy: 0.5238  
Epoch 23/800  
168/168 [=====] - 1s 5ms/step - loss: 1.4204 - accuracy: 0.5123 - val\_loss: 1.1664 - val\_accuracy: 0.5781  
Epoch 24/800  
168/168 [=====] - 1s 4ms/step - loss: 1.3674 - accuracy: 0.5268 - val\_loss: 1.0499 - val\_accuracy: 0.6548  
Epoch 25/800  
168/168 [=====] - 1s 5ms/step - loss: 1.3174 - accuracy: 0.5463 - val\_loss: 0.9300 - val\_accuracy: 0.6882  
Epoch 26/800  
168/168 [=====] - 1s 5ms/step - loss: 1.2657 - accuracy: 0.5605 - val\_loss: 0.9048 - val\_accuracy: 0.6897  
Epoch 27/800  
168/168 [=====] - 1s 5ms/step - loss: 1.2483 - accuracy: 0.5711 - val\_loss: 0.9503 - val\_accuracy: 0.6763  
Epoch 28/800  
168/168 [=====] - 1s 5ms/step - loss: 1.1901 - accuracy: 0.5952 - val\_loss: 0.8530 - val\_accuracy: 0.7150  
Epoch 29/800  
168/168 [=====] - 1s 6ms/step - loss: 1.1623 - accuracy: 0.6044 - val\_loss: 0.8687 - val\_accuracy: 0.7173  
Epoch 30/800  
168/168 [=====] - 1s 5ms/step - loss: 1.1273 - accuracy: 0.6138 - val\_loss: 0.8598 - val\_accuracy: 0.7121  
Epoch 31/800  
168/168 [=====] - 1s 5ms/step - loss: 1.0953 - accuracy: 0.6300 - val\_loss: 0.8026 - val\_accuracy: 0.7277  
Epoch 32/800  
168/168 [=====] - 1s 5ms/step - loss: 1.0795 - accuracy: 0.6453 - val\_loss: 0.7402 - val\_accuracy: 0.7560  
Epoch 33/800  
168/168 [=====] - 1s 4ms/step - loss: 1.0098 - accuracy: 0.6628 - val\_loss: 0.7267 - val\_accuracy: 0.7485  
Epoch 34/800  
168/168 [=====] - 1s 4ms/step - loss: 1.0077 - accuracy: 0.6656 - val\_loss: 0.7047 - val\_accuracy: 0.7805  
Epoch 35/800  
168/168 [=====] - 1s 5ms/step - loss: 0.9852 - accuracy: 0.6782 - val\_loss: 0.7622 - val\_accuracy: 0.7411  
Epoch 36/800  
168/168 [=====] - 2s 11ms/step - loss: 0.9664 - accuracy: 0.6775 - val\_loss: 0.6947 - val\_accuracy: 0.7641  
Epoch 37/800  
168/168 [=====] - 1s 4ms/step - loss: 0.9167 - accuracy: 0.6936 - val\_loss: 0.6373 - val\_accuracy: 0.7932  
Epoch 38/800  
168/168 [=====] - 1s 4ms/step - loss: 0.9023 - accuracy: 0.7070 - val\_loss: 0.6056 - val\_accuracy: 0.8065

Epoch 39/800  
168/168 [=====] - 1s 5ms/step - loss: 0.8859 - accuracy: 0.7096 - val\_loss: 0.6947 - val\_accuracy: 0.7790  
Epoch 40/800  
168/168 [=====] - 1s 5ms/step - loss: 0.8911 - accuracy: 0.7063 - val\_loss: 0.7163 - val\_accuracy: 0.7626  
Epoch 41/800  
168/168 [=====] - 1s 5ms/step - loss: 0.8608 - accuracy: 0.7160 - val\_loss: 0.6201 - val\_accuracy: 0.7865  
Epoch 42/800  
168/168 [=====] - 1s 4ms/step - loss: 0.8586 - accuracy: 0.7184 - val\_loss: 0.7056 - val\_accuracy: 0.7612  
Epoch 43/800  
168/168 [=====] - 1s 5ms/step - loss: 0.8082 - accuracy: 0.7307 - val\_loss: 0.6105 - val\_accuracy: 0.7902  
Epoch 44/800  
168/168 [=====] - 1s 5ms/step - loss: 0.7918 - accuracy: 0.7394 - val\_loss: 0.6486 - val\_accuracy: 0.7783  
Epoch 45/800  
168/168 [=====] - 1s 5ms/step - loss: 0.7865 - accuracy: 0.7372 - val\_loss: 0.5711 - val\_accuracy: 0.8222  
Epoch 46/800  
168/168 [=====] - 1s 4ms/step - loss: 0.7584 - accuracy: 0.7493 - val\_loss: 0.4940 - val\_accuracy: 0.8341  
Epoch 47/800  
168/168 [=====] - 1s 4ms/step - loss: 0.7589 - accuracy: 0.7500 - val\_loss: 0.7121 - val\_accuracy: 0.7746  
Epoch 48/800  
168/168 [=====] - 1s 5ms/step - loss: 0.7588 - accuracy: 0.7467 - val\_loss: 0.5285 - val\_accuracy: 0.8326  
Epoch 49/800  
168/168 [=====] - 1s 4ms/step - loss: 0.7295 - accuracy: 0.7599 - val\_loss: 0.6696 - val\_accuracy: 0.7954  
Epoch 50/800  
168/168 [=====] - 1s 4ms/step - loss: 0.7224 - accuracy: 0.7636 - val\_loss: 0.5189 - val\_accuracy: 0.8348  
Epoch 51/800  
168/168 [=====] - 1s 4ms/step - loss: 0.7113 - accuracy: 0.7706 - val\_loss: 0.5149 - val\_accuracy: 0.8318  
Epoch 52/800  
168/168 [=====] - 1s 5ms/step - loss: 0.7040 - accuracy: 0.7822 - val\_loss: 0.4644 - val\_accuracy: 0.8467  
Epoch 53/800  
168/168 [=====] - 1s 4ms/step - loss: 0.6747 - accuracy: 0.7760 - val\_loss: 0.4953 - val\_accuracy: 0.8296  
Epoch 54/800  
168/168 [=====] - 1s 4ms/step - loss: 0.6708 - accuracy: 0.7859 - val\_loss: 0.4516 - val\_accuracy: 0.8519  
Epoch 55/800  
168/168 [=====] - 1s 5ms/step - loss: 0.6672 - accuracy: 0.7835 - val\_loss: 0.5157 - val\_accuracy: 0.8356  
Epoch 56/800  
168/168 [=====] - 1s 4ms/step - loss: 0.6797 - accuracy: 0.7738 - val\_loss: 0.5063 - val\_accuracy: 0.8259  
Epoch 57/800  
168/168 [=====] - 1s 4ms/step - loss: 0.6529 - accuracy: 0.7831 - val\_loss: 0.4404 - val\_accuracy: 0.8564  
Epoch 58/800  
168/168 [=====] - 1s 4ms/step - loss: 0.6453 - accuracy: 0.7898 - val\_loss: 0.6866 - val\_accuracy: 0.7723  
Epoch 59/800  
168/168 [=====] - 1s 5ms/step - loss: 0.6697 - accuracy: 0.7844 - val\_loss: 0.7575 - val\_accuracy: 0.7812  
Epoch 60/800  
168/168 [=====] - 1s 4ms/step - loss: 0.6371 - accuracy: 0.7967 - val\_loss: 0.5183 - val\_accuracy: 0.8318  
Epoch 61/800  
168/168 [=====] - 1s 4ms/step - loss: 0.6327 - accuracy: 0.7930 - val\_loss: 0.5575 - val\_accuracy: 0.8155  
Epoch 62/800  
168/168 [=====] - 1s 5ms/step - loss: 0.6313 - accuracy: 0.8012 - val\_loss: 0.6615 - val\_accuracy: 0.7999  
Epoch 63/800  
168/168 [=====] - 1s 5ms/step - loss: 0.5993 - accuracy: 0.8052 - val\_loss: 0.4122 - val\_accuracy: 0.8676  
Epoch 64/800  
168/168 [=====] - 1s 5ms/step - loss: 0.6065 - accuracy: 0.8092 - val\_loss: 0.4945 - val\_accuracy: 0.8348  
Epoch 65/800  
168/168 [=====] - 1s 5ms/step - loss: 0.5875 - accuracy: 0.8088 - val\_loss: 0.4826 - val\_accuracy: 0.8527  
Epoch 66/800  
168/168 [=====] - 1s 4ms/step - loss: 0.5786 - accuracy: 0.8106 - val\_loss: 0.4357 - val\_accuracy: 0.8430  
Epoch 67/800  
168/168 [=====] - 1s 4ms/step - loss: 0.6017 - accuracy: 0.8067 - val\_loss: 0.4859 - val\_accuracy: 0.8504  
Epoch 68/800  
168/168 [=====] - 2s 11ms/step - loss: 0.5824 - accuracy: 0.8203 - val\_loss: 0.4263 - val\_accuracy: 0.8668  
Epoch 69/800  
168/168 [=====] - 1s 5ms/step - loss: 0.5884 - accuracy: 0.8166 - val\_loss: 0.4160 - val\_accuracy: 0.8757  
Epoch 70/800  
168/168 [=====] - 1s 4ms/step - loss: 0.5780 - accuracy: 0.8140 - val\_loss: 0.5490 - val\_accuracy:

y: 0.8304  
Epoch 71/800  
168/168 [=====] - 1s 4ms/step - loss: 0.5634 - accuracy: 0.8235 - val\_loss: 0.3835 - val\_accuracy: 0.8735  
Epoch 72/800  
168/168 [=====] - 1s 5ms/step - loss: 0.5479 - accuracy: 0.8250 - val\_loss: 0.4911 - val\_accuracy: 0.8333  
y: 0.8579  
Epoch 73/800  
168/168 [=====] - 1s 4ms/step - loss: 0.5637 - accuracy: 0.8212 - val\_loss: 0.4285 - val\_accuracy: 0.8780  
y: 0.8817  
Epoch 75/800  
168/168 [=====] - 1s 4ms/step - loss: 0.5213 - accuracy: 0.8279 - val\_loss: 0.3639 - val\_accuracy: 0.8817  
y: 0.8542  
Epoch 77/800  
168/168 [=====] - 1s 5ms/step - loss: 0.5351 - accuracy: 0.8274 - val\_loss: 0.4318 - val\_accuracy: 0.8571  
y: 0.8467  
Epoch 79/800  
168/168 [=====] - 1s 5ms/step - loss: 0.5207 - accuracy: 0.8350 - val\_loss: 0.3710 - val\_accuracy: 0.8899  
y: 0.8862  
Epoch 81/800  
168/168 [=====] - 1s 5ms/step - loss: 0.5043 - accuracy: 0.8348 - val\_loss: 0.5231 - val\_accuracy: 0.8393  
y: 0.8757  
Epoch 83/800  
168/168 [=====] - 1s 4ms/step - loss: 0.5217 - accuracy: 0.8363 - val\_loss: 0.4040 - val\_accuracy: 0.8713  
y: 0.8824  
Epoch 84/800  
168/168 [=====] - 1s 5ms/step - loss: 0.5017 - accuracy: 0.8337 - val\_loss: 0.3528 - val\_accuracy: 0.8757  
y: 0.8661  
Epoch 86/800  
168/168 [=====] - 1s 4ms/step - loss: 0.4719 - accuracy: 0.8464 - val\_loss: 0.4284 - val\_accuracy: 0.8493  
y: 0.8936  
Epoch 88/800  
168/168 [=====] - 1s 5ms/step - loss: 0.4935 - accuracy: 0.8451 - val\_loss: 0.4223 - val\_accuracy: 0.8616  
y: 0.8780  
Epoch 90/800  
168/168 [=====] - 1s 4ms/step - loss: 0.4637 - accuracy: 0.8516 - val\_loss: 0.3933 - val\_accuracy: 0.8854  
y: 0.8817  
Epoch 92/800  
168/168 [=====] - 1s 4ms/step - loss: 0.4775 - accuracy: 0.8458 - val\_loss: 0.3839 - val\_accuracy: 0.8951  
y: 0.8839  
Epoch 93/800  
168/168 [=====] - 1s 4ms/step - loss: 0.4518 - accuracy: 0.8527 - val\_loss: 0.3944 - val\_accuracy: 0.8795  
y: 0.8624  
Epoch 95/800  
168/168 [=====] - 1s 5ms/step - loss: 0.4442 - accuracy: 0.8581 - val\_loss: 0.3973 - val\_accuracy: 0.8981  
y: 0.8981  
Epoch 97/800  
168/168 [=====] - 1s 5ms/step - loss: 0.4546 - accuracy: 0.8560 - val\_loss: 0.3516 - val\_accuracy: 0.8891  
y: 0.9003  
Epoch 100/800  
168/168 [=====] - 1s 5ms/step - loss: 0.4501 - accuracy: 0.8612 - val\_loss: 0.3379 - val\_accuracy: 0.8936  
y: 0.9122  
Epoch 102/800

y: 0.9568  
Epoch 736/800  
168/168 [=====] - 1s 4ms/step - loss: 0.1236 - accuracy: 0.9643 - val\_loss: 0.2033 - val\_accuracy: 0.9568  
y: 0.9568  
Epoch 737/800  
168/168 [=====] - 1s 5ms/step - loss: 0.1249 - accuracy: 0.9617 - val\_loss: 0.1741 - val\_accuracy: 0.9621  
y: 0.9621  
Epoch 738/800  
168/168 [=====] - 1s 5ms/step - loss: 0.1220 - accuracy: 0.9587 - val\_loss: 0.1600 - val\_accuracy: 0.9680  
y: 0.9680  
Epoch 739/800  
168/168 [=====] - 1s 5ms/step - loss: 0.1199 - accuracy: 0.9652 - val\_loss: 0.1839 - val\_accuracy: 0.9606  
y: 0.9606  
Epoch 740/800  
168/168 [=====] - 1s 5ms/step - loss: 0.1074 - accuracy: 0.9656 - val\_loss: 0.1776 - val\_accuracy: 0.9583  
y: 0.9583  
Epoch 741/800  
168/168 [=====] - 1s 4ms/step - loss: 0.1099 - accuracy: 0.9667 - val\_loss: 0.1719 - val\_accuracy: 0.9658  
y: 0.9658  
Epoch 742/800  
168/168 [=====] - 1s 5ms/step - loss: 0.1129 - accuracy: 0.9630 - val\_loss: 0.1990 - val\_accuracy: 0.9591  
y: 0.9591  
Epoch 743/800  
168/168 [=====] - 1s 5ms/step - loss: 0.1141 - accuracy: 0.9619 - val\_loss: 0.1714 - val\_accuracy: 0.9650  
y: 0.9650  
Epoch 744/800  
168/168 [=====] - 1s 4ms/step - loss: 0.1115 - accuracy: 0.9632 - val\_loss: 0.1672 - val\_accuracy: 0.9650  
y: 0.9650  
Epoch 745/800  
168/168 [=====] - 1s 5ms/step - loss: 0.1037 - accuracy: 0.9708 - val\_loss: 0.1957 - val\_accuracy: 0.9613  
y: 0.9613  
Epoch 746/800  
168/168 [=====] - 1s 5ms/step - loss: 0.1216 - accuracy: 0.9632 - val\_loss: 0.1642 - val\_accuracy: 0.9717  
y: 0.9717  
Epoch 747/800  
168/168 [=====] - 1s 4ms/step - loss: 0.1145 - accuracy: 0.9637 - val\_loss: 0.1797 - val\_accuracy: 0.9635  
y: 0.9635  
Epoch 748/800  
168/168 [=====] - 1s 4ms/step - loss: 0.1119 - accuracy: 0.9626 - val\_loss: 0.1813 - val\_accuracy: 0.9606  
y: 0.9606  
Epoch 749/800  
168/168 [=====] - 1s 5ms/step - loss: 0.1162 - accuracy: 0.9628 - val\_loss: 0.1847 - val\_accuracy: 0.9628  
y: 0.9628  
Epoch 750/800  
168/168 [=====] - 2s 11ms/step - loss: 0.1015 - accuracy: 0.9691 - val\_loss: 0.1769 - val\_accuracy: 0.9598  
y: 0.9598  
Epoch 751/800  
168/168 [=====] - 1s 5ms/step - loss: 0.1078 - accuracy: 0.9658 - val\_loss: 0.1775 - val\_accuracy: 0.9606  
y: 0.9606  
Epoch 752/800  
168/168 [=====] - 1s 5ms/step - loss: 0.0996 - accuracy: 0.9686 - val\_loss: 0.1903 - val\_accuracy: 0.9606  
y: 0.9606  
Epoch 753/800  
168/168 [=====] - 1s 4ms/step - loss: 0.1070 - accuracy: 0.9647 - val\_loss: 0.1845 - val\_accuracy: 0.9613  
y: 0.9613  
Epoch 754/800  
168/168 [=====] - 1s 5ms/step - loss: 0.1082 - accuracy: 0.9667 - val\_loss: 0.2023 - val\_accuracy: 0.9554  
y: 0.9554  
Epoch 755/800  
168/168 [=====] - 1s 5ms/step - loss: 0.1285 - accuracy: 0.9602 - val\_loss: 0.2007 - val\_accuracy: 0.9591  
y: 0.9591  
Epoch 756/800  
168/168 [=====] - 1s 5ms/step - loss: 0.1184 - accuracy: 0.9639 - val\_loss: 0.1842 - val\_accuracy: 0.9628  
y: 0.9628  
Epoch 757/800  
168/168 [=====] - 1s 4ms/step - loss: 0.1174 - accuracy: 0.9652 - val\_loss: 0.1937 - val\_accuracy: 0.9613  
y: 0.9613  
Epoch 758/800  
168/168 [=====] - 1s 5ms/step - loss: 0.0981 - accuracy: 0.9699 - val\_loss: 0.1672 - val\_accuracy: 0.9665  
y: 0.9665  
Epoch 759/800  
168/168 [=====] - 1s 4ms/step - loss: 0.0919 - accuracy: 0.9701 - val\_loss: 0.1839 - val\_accuracy: 0.9688  
y: 0.9688  
Epoch 760/800  
168/168 [=====] - 1s 4ms/step - loss: 0.0911 - accuracy: 0.9725 - val\_loss: 0.1850 - val\_accuracy: 0.9665  
y: 0.9665  
Epoch 761/800  
168/168 [=====] - 1s 4ms/step - loss: 0.1127 - accuracy: 0.9645 - val\_loss: 0.1840 - val\_accuracy: 0.9583  
y: 0.9583  
Epoch 762/800  
168/168 [=====] - 1s 5ms/step - loss: 0.1140 - accuracy: 0.9671 - val\_loss: 0.1894 - val\_accuracy: 0.9591  
y: 0.9591  
Epoch 763/800  
168/168 [=====] - 1s 5ms/step - loss: 0.1173 - accuracy: 0.9643 - val\_loss: 0.1713 - val\_accuracy: 0.9680  
y: 0.9680  
Epoch 764/800  
168/168 [=====] - 1s 5ms/step - loss: 0.1101 - accuracy: 0.9647 - val\_loss: 0.1814 - val\_accuracy: 0.9621  
y: 0.9621  
Epoch 765/800  
168/168 [=====] - 1s 4ms/step - loss: 0.1099 - accuracy: 0.9669 - val\_loss: 0.1781 - val\_accuracy: 0.9643  
y: 0.9643  
Epoch 766/800  
168/168 [=====] - 1s 5ms/step - loss: 0.0930 - accuracy: 0.9715 - val\_loss: 0.1735 - val\_accuracy: 0.9628  
y: 0.9628  
Epoch 767/800

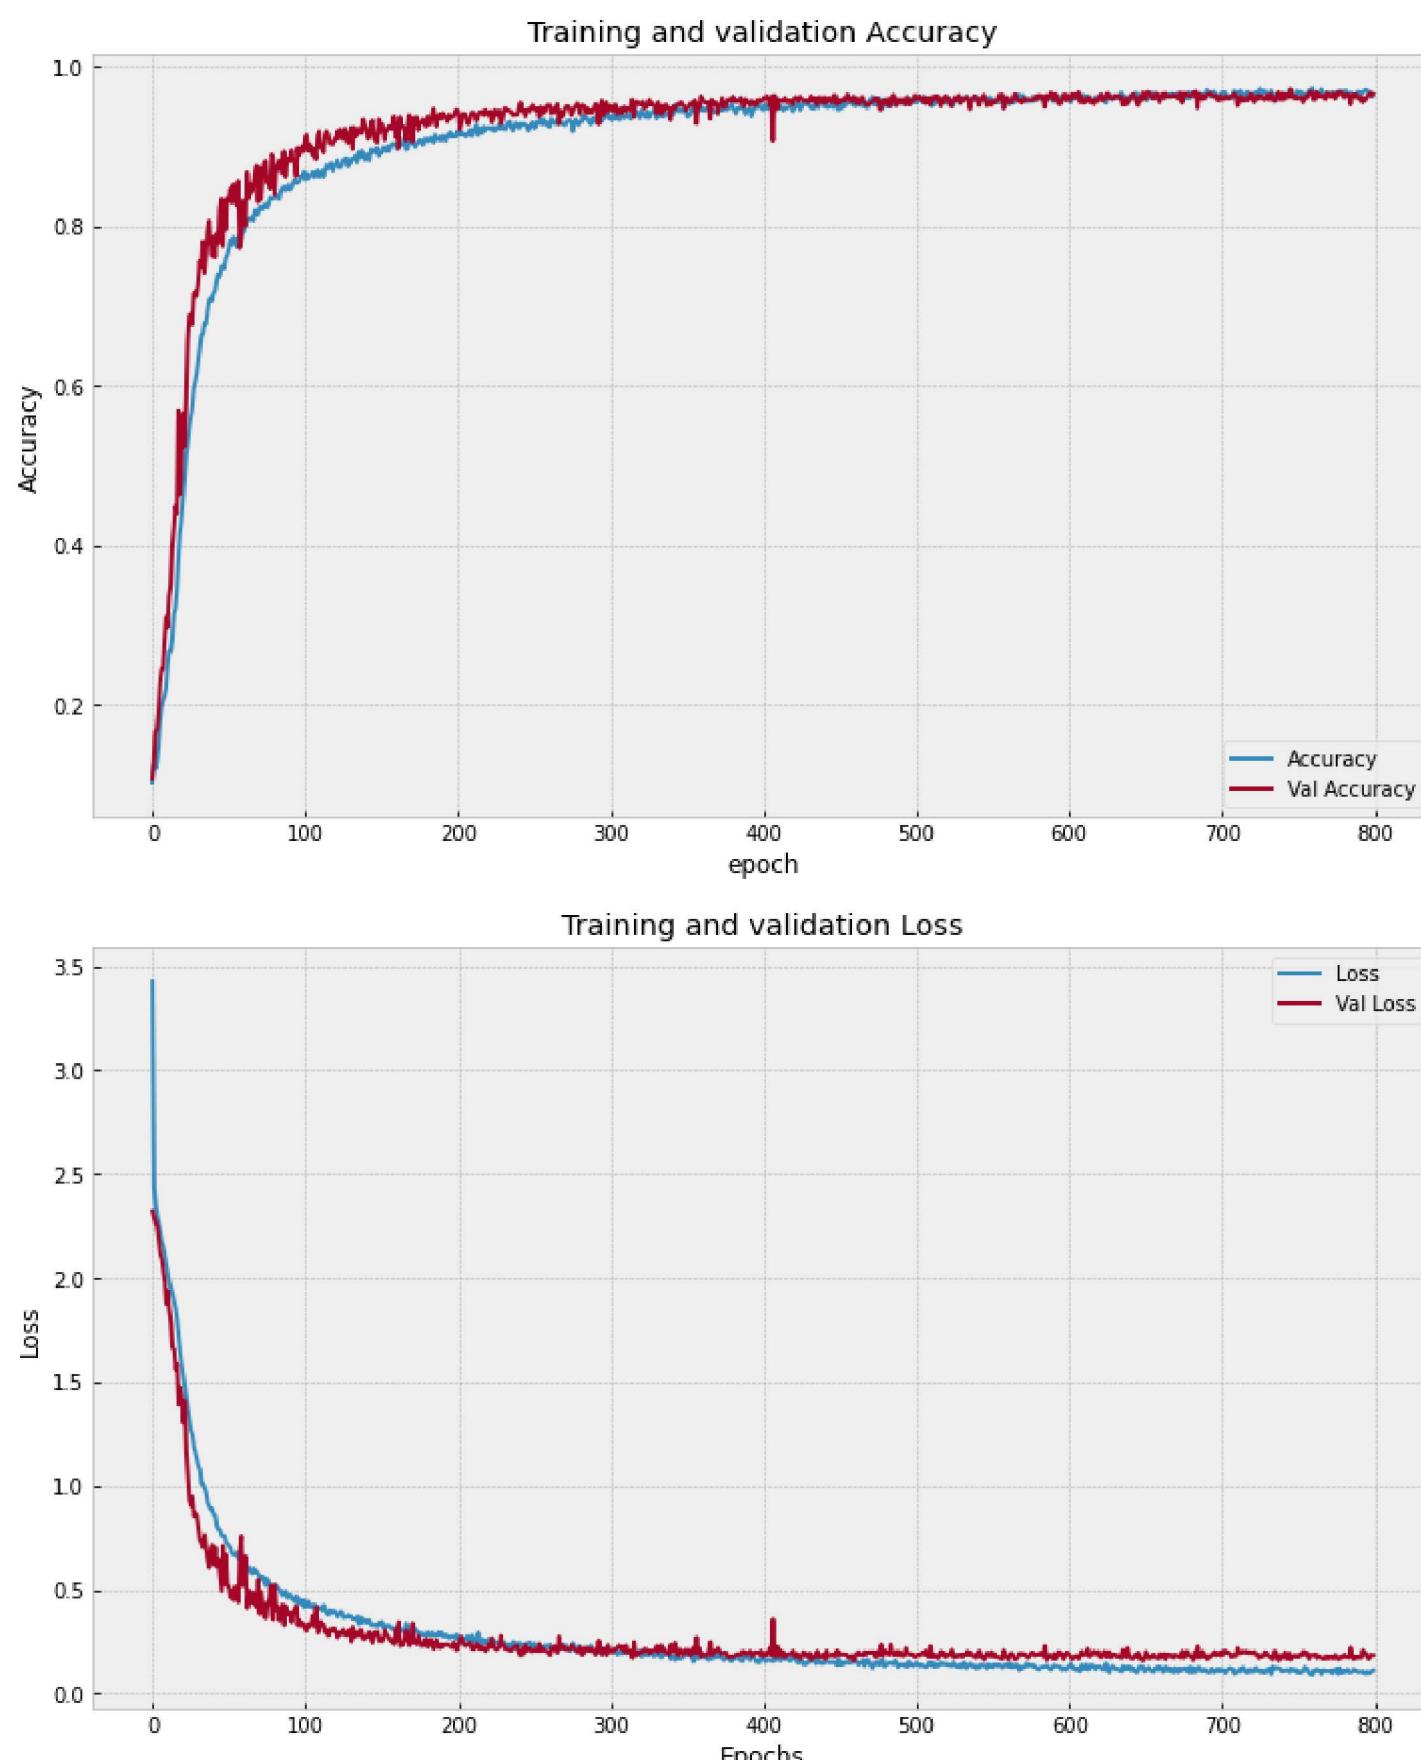
168/168 [=====] - 1s 5ms/step - loss: 0.1128 - accuracy: 0.9652 - val\_loss: 0.1678 - val\_accuracy: 0.9628  
Epoch 768/800  
168/168 [=====] - 1s 4ms/step - loss: 0.1131 - accuracy: 0.9660 - val\_loss: 0.1695 - val\_accuracy: 0.9606  
Epoch 769/800  
168/168 [=====] - 1s 5ms/step - loss: 0.1081 - accuracy: 0.9669 - val\_loss: 0.1804 - val\_accuracy: 0.9591  
Epoch 770/800  
168/168 [=====] - 1s 5ms/step - loss: 0.1099 - accuracy: 0.9661 - val\_loss: 0.1694 - val\_accuracy: 0.9673  
Epoch 771/800  
168/168 [=====] - 1s 5ms/step - loss: 0.1069 - accuracy: 0.9650 - val\_loss: 0.1771 - val\_accuracy: 0.9643  
Epoch 772/800  
168/168 [=====] - 1s 5ms/step - loss: 0.1106 - accuracy: 0.9641 - val\_loss: 0.1729 - val\_accuracy: 0.9628  
Epoch 773/800  
168/168 [=====] - 1s 5ms/step - loss: 0.0979 - accuracy: 0.9680 - val\_loss: 0.1780 - val\_accuracy: 0.9621  
Epoch 774/800  
168/168 [=====] - 1s 4ms/step - loss: 0.1145 - accuracy: 0.9645 - val\_loss: 0.1693 - val\_accuracy: 0.9665  
Epoch 775/800  
168/168 [=====] - 1s 4ms/step - loss: 0.1068 - accuracy: 0.9665 - val\_loss: 0.1857 - val\_accuracy: 0.9598  
Epoch 776/800  
168/168 [=====] - 1s 4ms/step - loss: 0.1205 - accuracy: 0.9643 - val\_loss: 0.1787 - val\_accuracy: 0.9621  
Epoch 777/800  
168/168 [=====] - 1s 5ms/step - loss: 0.0953 - accuracy: 0.9691 - val\_loss: 0.1652 - val\_accuracy: 0.9680  
Epoch 778/800  
168/168 [=====] - 1s 5ms/step - loss: 0.1229 - accuracy: 0.9652 - val\_loss: 0.1703 - val\_accuracy: 0.9665  
Epoch 779/800  
168/168 [=====] - 1s 5ms/step - loss: 0.0991 - accuracy: 0.9695 - val\_loss: 0.1778 - val\_accuracy: 0.9680  
Epoch 780/800  
168/168 [=====] - 1s 5ms/step - loss: 0.1097 - accuracy: 0.9684 - val\_loss: 0.1746 - val\_accuracy: 0.9650  
Epoch 781/800  
168/168 [=====] - 1s 4ms/step - loss: 0.1034 - accuracy: 0.9673 - val\_loss: 0.1868 - val\_accuracy: 0.9613  
Epoch 782/800  
168/168 [=====] - 1s 5ms/step - loss: 0.1080 - accuracy: 0.9699 - val\_loss: 0.1757 - val\_accuracy: 0.9643  
Epoch 783/800  
168/168 [=====] - 1s 5ms/step - loss: 0.1167 - accuracy: 0.9656 - val\_loss: 0.1784 - val\_accuracy: 0.9650  
Epoch 784/800  
168/168 [=====] - 1s 4ms/step - loss: 0.1016 - accuracy: 0.9663 - val\_loss: 0.1741 - val\_accuracy: 0.9643  
Epoch 785/800  
168/168 [=====] - 1s 5ms/step - loss: 0.1031 - accuracy: 0.9669 - val\_loss: 0.2221 - val\_accuracy: 0.9524  
Epoch 786/800  
168/168 [=====] - 1s 5ms/step - loss: 0.1051 - accuracy: 0.9665 - val\_loss: 0.1816 - val\_accuracy: 0.9635  
Epoch 787/800  
168/168 [=====] - 1s 4ms/step - loss: 0.1177 - accuracy: 0.9656 - val\_loss: 0.1666 - val\_accuracy: 0.9650  
Epoch 788/800  
168/168 [=====] - 1s 5ms/step - loss: 0.1124 - accuracy: 0.9654 - val\_loss: 0.1779 - val\_accuracy: 0.9613  
Epoch 789/800  
168/168 [=====] - 2s 11ms/step - loss: 0.1123 - accuracy: 0.9650 - val\_loss: 0.1721 - val\_accuracy: 0.9650  
Epoch 790/800  
168/168 [=====] - 1s 5ms/step - loss: 0.1036 - accuracy: 0.9684 - val\_loss: 0.1779 - val\_accuracy: 0.9613  
Epoch 791/800  
168/168 [=====] - 1s 5ms/step - loss: 0.1049 - accuracy: 0.9680 - val\_loss: 0.1733 - val\_accuracy: 0.9591  
Epoch 792/800  
168/168 [=====] - 1s 5ms/step - loss: 0.1062 - accuracy: 0.9656 - val\_loss: 0.1732 - val\_accuracy: 0.9621  
Epoch 793/800  
168/168 [=====] - 1s 4ms/step - loss: 0.1087 - accuracy: 0.9682 - val\_loss: 0.2100 - val\_accuracy: 0.9554  
Epoch 794/800  
168/168 [=====] - 1s 5ms/step - loss: 0.0984 - accuracy: 0.9693 - val\_loss: 0.1946 - val\_accuracy: 0.9591  
Epoch 795/800  
168/168 [=====] - 1s 4ms/step - loss: 0.0962 - accuracy: 0.9708 - val\_loss: 0.1935 - val\_accuracy: 0.9583  
Epoch 796/800  
168/168 [=====] - 1s 5ms/step - loss: 0.1076 - accuracy: 0.9680 - val\_loss: 0.1726 - val\_accuracy: 0.9621  
Epoch 797/800  
168/168 [=====] - 1s 4ms/step - loss: 0.0982 - accuracy: 0.9688 - val\_loss: 0.1669 - val\_accuracy: 0.9643  
Epoch 798/800  
168/168 [=====] - 1s 4ms/step - loss: 0.1003 - accuracy: 0.9680 - val\_loss: 0.1867 - val\_accuracy: 0.9665

```

Epoch 799/800
168/168 [=====] - 1s 4ms/step - loss: 0.1096 - accuracy: 0.9637 - val_loss: 0.1824 - val_accuracy: 0.9628
Epoch 800/800
168/168 [=====] - 1s 5ms/step - loss: 0.1129 - accuracy: 0.9639 - val_loss: 0.1847 - val_accuracy: 0.9658
Model Fitted

```

In [10]: `summarize_diagnostics(history)`



We can see that we are getting a flat line around 600-800 epochs with accuracy greater than 95%, and both Validation and Training accuracy and loss are close, so there is no underfitting and overfitting, but val\_loss is increasing after some 700 epochs, so we should use epochs between 600 and 700. So, this model is good, so we can train our model with these parameters on complete data\_train set.

**Below are the final functions used to Load, Preprocess amd Train our full data.**

In [12]: `def final_training_complete(dataX, dataY):
 model = define_model(learn_rate = 0.0007)
 trainX, trainY = dataX, dataY

 # fit model
 history = model.fit(trainX, trainY, epochs = 650, batch_size = 32, verbose=0)
 print("Model Fitted")
 return model, history`

In [13]: `def load_full_dataset():
 # Load dataset
 data_rgb = []
 kernel = np.ones((4,4),np.uint8)
 for i in range(train_data_length):
 data_rgb.append(data_train[:,i].reshape(300,300))
 data_rgb[i] = cv2.medianBlur(data_rgb[i], 3)
 data_rgb[i] = cv2.morphologyEx(data_rgb[i], cv2.MORPH_OPEN, kernel)
 data_rgb[i] = cv2.resize(data_rgb[i], (50,50), interpolation=cv2.INTER_AREA)
 data_rgb = np.array(data_rgb)
 print(data_rgb.shape)`

```
# reshape dataset to have a single channel
data_rgb = data_rgb.reshape((data_rgb.shape[0], 50,50, 1))

# one hot encode target values
labels = to_categorical(labels_train)
print(data_rgb.shape, labels.shape)
return data_rgb, labels
```

In [14]:

```
def prep_final_pixels(data):
    # convert from integers to floats
    data_norm = data.astype('float32')

    # he_uniformize to range 0-1
    data_norm = (data_norm) / 255.0

    # return he_uniformized images
    return data_norm
```

In this cell all final functions are called to train our dataset on all images.

In [15]:

```
data, target = load_full_dataset()

# prepare pixel data
data = prep_final_pixels(data)

# training the model
model, history = final_training_complete(data, target)
```

(6720, 50, 50)  
(6720, 50, 50, 1) (6720, 10)  
Model Fitted

In [16]:

```
plt.figure(figsize=(12,7))
plt.plot(history.history['accuracy'])
plt.plot(history.history['loss'])
plt.title('Training accuracy and Loss')
plt.ylabel('Accuracy and Loss')
plt.xlabel('epoch')
plt.legend(['Accuracy', 'Loss'], loc='center right')
plt.show()
```



In [17]:

```
# model.save("Final_CNN")
```

2022-04-20 14:16:39.478446: W tensorflow/python/util/util.cc:348] Sets are not currently considered sequences, but this may change in the future, so consider avoiding using them.  
INFO:tensorflow:Assets written to: Final\_CNN/assets

Finally, we saved our model as "Final\_CNN" and will use this trained model in our Easy\_Test and Hard\_Test files.